

Banned Books — Open Censorship Core: Data Descriptor

- **Dataset:** Banned Books — Open Censorship Core
 - **Publisher:** Banned Books (<https://www.banned-books.org>)
 - **Creator:** Ludo Raedts
 - **ORCID:** <https://orcid.org/0009-0006-8358-7119>
 - **License:** CC-BY-4.0 (<https://creativecommons.org/licenses/by/4.0/>)
 - **Concept DOI:** [10.5281/zenodo.20511553](https://doi.org/10.5281/zenodo.20511553) (version-independent — always resolves to the latest version).
 - **Companion files:** `books.csv`, `authors.csv`, `bans.csv`, `ban_reasons.csv`, `ban_sources.csv`, `countries.csv`, `schema.json`, `README.md`, `LICENSE.txt`.
-

1. Purpose & scope

Banned Books is a catalogue of books that have been banned, restricted, or sustainably challenged anywhere in the world, with a source citation behind every entry. This deposit is the **open, citeable core** of that catalogue: the structured, verifiable facts — who banned what, where, when, why, and on whose authority — together with the reason taxonomy and the source citations that let a researcher check each record.

It is deliberately **not** identical to the commercial dataset sold at [/dataset](#). The split is principled:

- **Open (this deposit, CC-BY-4.0):** facts about censorship and the citations that ground them.
- **Commercial (separate license):** editorial prose (book and ban descriptions, censorship context), enrichment (ISBNs, cover images, author biographies and photos, edition data), and convenience formats (a single denormalised JSON file and a ready-to-query SQLite database).

The reason **taxonomy** (the reason *slug*, e.g. `lgbtq`, `political`) is open; the written description paragraph that interprets a ban is commercial.

Coverage at the snapshot used to write this descriptor (figures move as the catalogue grows — the deposited files are the authority):

- ~14,000 books, ~8,650 authors
- ~28,730 ban events across 119 countries (including defunct states: USSR, East Germany, Czechoslovakia, Yugoslavia)
- ~52,800 ban–reason links and ~28,800 ban–source citations

Scope. This deposit covers *books*. Records must describe a real removal/restriction with an institutional actor and a documented decision (see §4). The catalogue does not carry a media-type field, so non-book media are not

systematically distinguished or excluded at the schema level; the editorial intent is books, but the data offers no hard guarantee that every row is one.

2. Data model

Six CSV tables. All join keys are public, human-readable slugs or codes; the only surrogate key is `ban_id` (a row identifier needed because a ban has no natural slug). `schema.json` holds the machine-readable column types.




Table	Grain	Primary key	Foreign keys
books.csv	one row per work	slug	author_slugs (pipe-separated) → authors.slug
authors.csv	one row per author	slug	—
bans.csv	one row per ban event	ban_id	book_slug → books.slug; country_code → countries.code
ban_reasons.csv	many rows per ban	(ban_id,reason_slug)	ban_id → bans.ban_id
ban_sources.csv	many rows per ban	—	ban_id → bans.ban_id
countries.csv	one row per country	code	—

Join keys.

- books.author_slugs is a pipe-separated (|) list — split it, then join each element to authors.slug. This carries the book↔author relationship without a separate join table.
- bans.book_slug → books.slug; bans.country_code → countries.code.
- ban_reasons.ban_id and ban_sources.ban_id → bans.ban_id.

Conventions. Array fields are pipe-separated (|). NULLs are encoded as the empty string. Years are integers and may be negative or pre-1000 CE for ancient works.

3. Field definitions

Authoritative, machine-readable types live in `schema.json`. Summary:

books.csv — slug (PK), title (canonical published title), first_published_year (int, nullable), original_language (code, nullable), author_slugs (pipe-separated authors.slug, nullable).

`authors.csv` — `slug` (PK), `display_name` (slug-canonical, Anglo-friendly form), `birth_country` (code, nullable), `is_placeholder` (boolean: `true` = aggregate / non-attributable bucket entry such as "Anonymous" or "Various Authors" that groups unrelated works). *No biographies* — those are commercial.

`bans.csv` — `ban_id` (PK, surrogate), `book_slug` (FK), `country_code` (FK), `year_started` / `year_ended` (int, nullable; `year_ended` empty if in force/unknown), `action_type` (one of banned, restricted, challenged), `status` (one of active, historical, rescinded), `scope` (taxonomy slug: school, government, prison, ...; empty if unscoped). `banned` and `active` dominate; `challenged` and `rescinded` are smaller categories.

Withheld rows. A tiny number of bans with an indeterminate status (`status` = 'unclear'; currently 2 rows) exist in the live catalogue but are deliberately omitted from this open export. As a result the open ban count is marginally lower than the headline figure on *banned-books.org*. The DB is not modified — these rows simply aren't deposited.

`ban_reasons.csv` — `ban_id` (FK), `reason_slug` (stable taxonomy slug), `reason_label` (English label). Zero or more rows per ban.

`ban_sources.csv` — `ban_id` (FK), `source_name`, `source_url` (nullable), `source_type` (nullable), `verification_status` (`verified` = URL works and archived; `pending` = archive attempt failed; `broken` = URL 4xx/5xx; `unverified` = never attempted; empty on pre-pipeline rows), `accessed_at` (ISO date, nullable), `locator` (in-source page/row/entry id, nullable). Zero or more rows per ban.

`countries.csv` — `code` (PK; ISO 3166-1 alpha-2 where applicable, custom codes for defunct states), `name_en`.

4. Methodology — how bans are sourced and verified

Data is drawn from PEN America, the American Library Association (ALA), Wikipedia, Reporters Without Borders, Index on Censorship, government records, and national-press reporting. Every ban entry links to its source via `ban_sources.csv`.

What counts as a ban. A broad definition: any formal removal, import restriction, publication prohibition, or sustained challenge that results in a book becoming unavailable through an official channel. Both *hard* bans (legally prohibited nationwide) and *soft* bans (withdrawn from a school or library after pressure) are included, provided there is a documented decision with an institutional actor.

What is excluded. Books that are merely out of print, books a publisher chose not to distribute, and books that are culturally stigmatised but legally available. The restriction must have an institutional actor and a documented decision.

Verification. Sources carry a `verification_status`. Archive-verification is an ongoing process and the field is mostly unpopulated so far. At the snapshot used here, of ~820 source rows the distribution is roughly: `unverified` ≈ 89%, `empty/(null)` ≈ 10%, `verified` = 5 rows, `pending` = 2 rows, and `broken` = 0.

In practical terms, filtering on `verification_status = 'verified'` currently returns a near-empty subset (a handful of rows) and is not yet a usable quality gate. Treat it as a lever that becomes useful as verification runs progress, not as something to filter on today. The catalogue is built mostly by automated import from public sources and then enriched and reviewed; not every entry has been individually checked by a human.

Per-record data quality. A composite per-record `data_quality_status` (`confident` / `default` / `flagged`) exists in the underlying catalogue, but it is intentionally **not** included in this open release; it is available in the commercial dataset. The open data does not provide an equivalent quality signal — the source citations in `ban_sources.csv` let you inspect provenance, but they are not a substitute for the composite status.

5. How to count — distinct books vs. raw events

Rank on distinct books or distinct countries, never on raw ban-row counts.

`bans.csv` holds **one row per event**. A US title removed across many school districts produces many rows; a book banned nationwide by a single government decision appears once. PEN America's per-district granularity therefore inflates the United States' raw row count roughly 2–3× relative to distinct titles.

- **Canonical metric for "how widely banned":** count **distinct** `book_slug` (and/or **distinct** `country_code`).
- **Raw `bans.csv` row counts** are supporting detail only — never a repression ranking.

Worked example (snapshot figures): ~28,700 raw ban rows resolve to ~14,000 distinct books — the gap is overwhelmingly US school-district events for a smaller set of titles. A country ranking built on raw rows would massively overstate the US relative to a ranking built on distinct banned titles.

The export is intentionally **not pre-aggregated**: every event is preserved so researchers can choose their own grain. This section exists so they don't mistake the event grain for a title count.

Status semantics. `rescinded` is a genuine category, not a synonym of `historical`: it marks a ban that was formally *lifted in a later year* (nearly all such rows — 51 of 52 at the current snapshot — carry a `year_ended`), whereas `historical` simply means no longer in force without asserting a documented reversal.

6. Known gaps and limitations

This dataset is honest about what it is not.

1. **Documentation bias — the US "counts more than it bans more."** The United States dominates the catalogue not because American libraries censor the most books, but because America *counts* them: PEN America and the ALA systematically record school-district removals and challenges, working with local journalists and librarians. No comparable reporting infrastructure exists elsewhere. A high US count reflects a watchful civil society; a low count for an authoritarian state reflects the absence of one.
2. **Authoritarian under-coverage.** Iran, China, North Korea, Saudi Arabia, and Belarus maintain comprehensive censorship regimes in which whole categories of literature are simply unavailable — but with no free press or civil-society watchdog, those bans rarely surface in any accessible database. The catalogue records only a fraction of what likely occurs in closed societies. **Absence of records is not evidence of freedom.**
3. **English-language reporting skew.** Sources we can index are disproportionately English-language, biasing coverage toward the Anglophone world and well-reported events.
4. **School-bans are structurally different.** A school-board removal is a local administrative decision, not national law; the book remains available elsewhere. The catalogue records both school and government bans but distinguishes them via `scope` and `status` — do not treat them as equivalent.
5. **Language / classification subset.** A subset of records carries imperfect `original_language` classification (automated language inference is error-prone for short titles and transliterated works). Treat `original_language` as indicative, not authoritative, for fine-grained linguistic analysis.
6. **Automated import, partial human review.** Most records originate from automated pipelines. Broad strokes (title, author, ban country) are reliable; narrower details may be provisional. `ban_sources.verification_status` is mostly `unverified` today (see §4), so it cannot yet stand in for record-level quality filtering.
7. **Placeholder / aggregate author entries.** A small number of `authors` rows are non-attributable buckets ("Anonymous", "Various Authors") that group unrelated works. They are flagged `is_placeholder = true`; filter them out before computing per-author statistics so a bucket isn't treated as a single author.

7. Licensing

This deposit — the open censorship core — is released under the **Creative Commons Attribution 4.0 International License (CC-BY-4.0)**. You may share and adapt the data for any purpose, including commercially, provided you give attribution (see §8).

Not included in this open release (available under a separate commercial license at <https://www.banned-books.org/dataset>):

- Editorial prose: book descriptions, ban descriptions, censorship context, extended context.
- Enrichment: ISBN-13, cover images and cover status, author biographies and photos, bookshop/edition data.
- Convenience formats: the single-file denormalised JSON and the SQLite database.

The principle: **facts about censorship are open; editorial prose and convenience formats are paid.**

8. Citation guidance

Cite the **concept DOI** — the version-independent Zenodo DOI that always resolves to the latest version — not a specific version DOI, unless you need to pin an exact snapshot for reproducibility.

Raedts, Ludo. Banned Books — Open Censorship Core. banned-books.org. Zenodo. CC-BY-4.0. ORCID: <https://orcid.org/0009-0006-8358-7119>. DOI: 10.5281/zenodo.20511553.

When attributing in prose or visualisations: **"Banned Books (banned-books.org), CC-BY-4.0"**, with a link to the DOI. If you publish a ranking or aggregate, please state whether it is built on distinct books / countries or on raw ban events (§5), so readers can interpret it correctly.